

Data frames, matrices, & subsetting

Lecture 05

Dr. Colin Rundel

Matrices and Arrays

Matrices

R supports the creation of 2d data structures (rows and columns) of atomic vector types.

Generally these are formed via a call to `matrix()`.

```
1 matrix(1:4, nrow=2, ncol=2)
```

```
 [,1] [,2]  
[1,] 1 3  
[2,] 2 4
```

```
1 matrix(c(TRUE, FALSE), 2, 2)
```

```
 [,1] [,2]  
[1,] TRUE TRUE  
[2,] FALSE FALSE
```

```
1 matrix(LETTERS[1:6], 2)
```

```
 [,1] [,2] [,3]  
[1,] "A" "C" "E"  
[2,] "B" "D" "F"
```

```
1 matrix(6:1 / 2, ncol = 2)
```

```
 [,1] [,2]  
[1,] 3.0 1.5  
[2,] 2.5 1.0  
[3,] 2.0 0.5
```

Data ordering

Matrices in R use column major ordering (data is sorted in column order not row order).

```
1 (m = matrix(1:6, nrow=2, ncol=3))
```

```
 [,1] [,2] [,3]  
[1,] 1 3 5  
[2,] 2 4 6
```

```
1 c(m)
```

```
[1] 1 2 3 4 5 6
```

```
1 (n = matrix(1:6, nrow=2, ncol=3))
```

```
 [,1] [,2] [,3]  
[1,] 1 3 5  
[2,] 2 4 6
```

```
1 c(n)
```

```
[1] 1 2 3 4 5 6
```

When creating the matrix we can populate the matrix via row, but that data will still be stored in column major order.

```
1 (x = matrix(1:6, nrow=2, ncol=3, byrow = TRUE))
```

```
 [,1] [,2] [,3]  
[1,] 1 2 3  
[2,] 4 5 6
```

```
1 c(x)
```

```
[1] 1 4 2 5 3 6
```

```
1 (y = matrix(1:6, nrow=3, ncol=2, byrow=TRUE))
```

```
 [,1] [,2]  
[1,] 1 2  
[2,] 3 4  
[3,] 5 6
```

```
1 c(y)
```

```
[1] 1 3 5 2 4 6
```

Matrix structure

```
1 m = matrix(1:4, ncol=2, nrow=2)
```

```
1 typeof(m)
```

```
[1] "integer"
```

```
1 mode(m)
```

```
[1] "numeric"
```

```
1 class(m)
```

```
[1] "matrix" "array"
```

```
1 attributes(m)
```

```
$dim
```

```
[1] 2 2
```

Matrices (and arrays) are just atomic vectors with a `dim` attribute attached (they do not have a `class` attribute, but they do have a `class`).

```
1 n = letters[1:6]
2 dim(n) = c(2L, 3L)
3
4 n
```

```
 [,1] [,2] [,3]
[1,] "a"   "c"   "e"
[2,] "b"   "d"   "f"
```

```
1 class(n)
```

```
[1] "matrix" "array"
```

```
1 o = letters[1:6]
2 attr(o, "dim") = c(2L, 3L)
3
4 o
```

```
 [,1] [,2] [,3]
[1,] "a"   "c"   "e"
[2,] "b"   "d"   "f"
```

```
1 class(o)
```

```
[1] "matrix" "array"
```

Arrays

Arrays are just an n-dimensional extension of matrices and are defined by adding the appropriate dimension sizes.

```
1 array(1:8, dim = c(2,2,2))
```

```
, , 1
```

```
[,1] [,2]  
[1,] 1 3  
[2,] 2 4
```

```
, , 2
```

```
[,1] [,2]  
[1,] 5 7  
[2,] 6 8
```

```
1 array(letters[1:6], dim = c(2,1,3))
```

```
, , 1
```

```
[,1]  
[1,] "a"  
[2,] "b"
```

```
, , 2
```

```
[,1]  
[1,] "c"  
[2,] "d"
```

```
, , 3
```

```
[,1]  
[1,] "e"  
[2,] "f"
```

Data Frames

Data Frames

A data frame is how R handles heterogeneous tabular data (i.e. a table of rows and columns) and is one of the most commonly used data structure in R.

```
1 df = data.frame(  
2   x = 1:3,  
3   y = c("a", "b", "c"),  
4   z = c(TRUE)  
5 )
```

```
x y     z  
1 1 a TRUE  
2 2 b TRUE  
3 3 c TRUE
```

R represents data frames using a *list* of equal length *vectors*.

```
1 str(df)
```



```
'data.frame': 3 obs. of 3 variables:  
 $ x: int 1 2 3  
 $ y: chr "a" "b" "c"  
 $ z: logi TRUE TRUE TRUE
```

Data Frame Structure

```
1 typeof(df)
```

```
[1] "list"
```

```
1 class(df)
```

```
[1] "data.frame"
```

```
1 attributes(df)
```

```
$names
```

```
[1] "x" "y" "z"
```

```
$class
```

```
[1] "data.frame"
```

```
$row.names
```

```
[1] 1 2 3
```

```
1 str(unclass(df))
```

```
List of 3
```

```
$ x: int [1:3] 1 2 3
```

```
$ y: chr [1:3] "a" "b" "c"
```

```
$ z: logi [1:3] TRUE TRUE TRUE
```

```
- attr(*, "row.names")= int [1:3] 1 2
```

```
3
```

Build your own data.frame

```
1 df = list(x = 1:3, y = c("a", "b", "c"), z = c(TRUE, TRUE, TRUE))
```

```
1 attr(df,"class") = "data.frame"
2 df
```

```
[1] x y z
<0 rows> (or 0-length row.names)
```

```
1 attr(df,"row.names") = 1:3
2 df
```

	x	y	z
1	1	a	TRUE
2	2	b	TRUE
3	3	c	TRUE

```
1 str(df)
```

```
'data.frame': 3 obs. of 3 variables:
$ x: int 1 2 3
$ y: chr "a" "b" "c"
$ z: logi TRUE TRUE TRUE
```

```
1 is.data.frame(df)
```

```
[1] TRUE
```

Strings (Characters) vs Factors

Previous to R v4.0, the default behavior of data frames was to convert character data into factors. Sometimes this was useful, but mostly it wasn't.

It is important to know what type/class you are working with. This behavior can be changed using the `stringsAsFactors` argument to `data.frame` and related functions (e.g. `read.csv`, `read.table`, etc.).

```
1 df = data.frame(x = 1:3, y = c("a", "b", "c"),
                  stringsAsFactors = TRUE)
2
3 df
```

```
x y
1 1 a
2 2 b
3 3 c
```

```
1 str(df)
```

```
'data.frame': 3 obs. of 2 variables:
$ x: int 1 2 3
$ y: Factor w/ 3 levels "a","b","c": 1 2 3
```

```
1 df = data.frame(x = 1:3, y = c("a", "b", "c"),
                  stringsAsFactors = FALSE)
2
3 df
```

```
x y
1 1 a
2 2 b
3 3 c
```

```
1 str(df)
```

```
'data.frame': 3 obs. of 2 variables:
$ x: int 1 2 3
$ y: chr "a" "b" "c"
```

Length Coercion

For data frames on creation the lengths of the component vectors will be coerced to match, however if they not multiples then there will be an error (previously this produced a warning).

```
1 data.frame(x = 1:3, y = c("a"))
```

```
x y
1 1 a
2 2 a
3 3 a
```

```
1 data.frame(x = 1:3, y = c("a", "b"))
```

Error in data.frame(x = 1:3, y = c("a", "b")): arguments imply differing number of rows: 3, 2

```
1 data.frame(x = 1:3, y = character())
```

Error in data.frame(x = 1:3, y = character()): arguments imply differing number of rows: 3, 0

Subsetting

Subsetting in General

R has three subsetting operators (`[`, `[[`, and `$`). The behavior of these operators will depend on the object (class) they are being used with.

In general there are 6 different types of subsetting that can be performed:

- Positive integer
- Negative integer
- Logical value
- Empty / NULL
- Zero valued
- Character value (names)

Positive Integer subsetting

Returns elements at the given location(s)

```
1 x = c(1,4,7)
2 y = list(1,4,7)
```

```
1 x[1]
```

```
[1] 1
```

```
1 x[c(1,3)]
```

```
[1] 1 7
```

```
1 x[c(1,1)]
```

```
[1] 1 1
```

```
1 x[c(1.9,2.1)]
```

```
[1] 1 4
```

```
1 str( y[1] )
```

List of 1

\$: num 1

```
1 str( y[c(1,3)] )
```

List of 2

\$: num 1

\$: num 7

```
1 str( y[c(1,1)] )
```

List of 2

\$: num 1

\$: num 1

```
1 str( y[c(1.9,2.1)] )
```

List of 2

\$: num 1

\$: num 4

Negative Integer subsetting

Excludes elements at the given location(s)

```
1 x = c(1,4,7)
```

```
1 x[-1]
```

```
[1] 4 7
```

```
1 x[-c(1,3)]
```

```
[1] 4
```

```
1 x[c(-1,-1)]
```

```
[1] 4 7
```

```
1 x[c(-1,2)]
```

Error in x[c(-1, 2)]: only 0's may be mixed with negative subscripts

```
1 y[c(-1,2)]
```

Error in y[c(-1, 2)]: only 0's may be mixed with negative subscripts

```
1 y = list(1,4,7)
```

```
1 str( y[-1] )
```

List of 2

\$: num 4

\$: num 7

```
1 str( y[-c(1,3)] )
```

List of 1

\$: num 4

Logical Value Subsetting

Returns elements that correspond to `TRUE` in the logical vector. Length of the logical vector is coerced to be the same as the vector being subsetted.

```
1 x = c(1, 4, 7, 12)
```

```
1 x[c(TRUE, TRUE, FALSE, TRUE)]
```

```
[1] 1 4 12
```

```
1 x[c(TRUE, FALSE)]
```

```
[1] 1 7
```

```
1 x[x %% 2 == 0]
```

```
[1] 4 12
```

```
1 y = list(1, 4, 7, 12)
```

```
1 str( y[c(TRUE, TRUE, FALSE, TRUE)] )
```

List of 3

```
$ : num 1  
$ : num 4  
$ : num 12
```

```
1 str( y[c(TRUE, FALSE)] )
```

List of 2

```
$ : num 1  
$ : num 7
```

```
1 str( y[y %% 2 == 0] )
```

Error in y%%2: non-numeric argument to binary operator

Empty Subsetting

Returns the original vector, this is not the same as subsetting with `NULL`

```
1 x = c(1,4,7)
```

```
1 x[ ]
```

```
[1] 1 4 7
```

```
1 x[NULL]
```

```
numeric(0)
```

```
1 y = list(1,4,7)
```

```
1 str(y[ ])
```

List of 3

```
$ : num 1  
$ : num 4  
$ : num 7
```

```
1 str(y[NULL])
```

`list()`

Zero subsetting

Returns an empty vector (of the same type), this is the same as subsetting with `NULL`

```
1 x = c(1, 4, 7)
```

```
1 x[0]
```

```
numeric(0)
```

```
1 y = list(1, 4, 7)
2 str(y[0])
```

```
list()
```

0s can be mixed with either positive or negative integers for subsetting (they are ignored)

```
1 x[c(0, 1)]
```

```
[1] 1
```

```
1 y[c(0, 1)]
```

```
[[1]]
```

```
[1] 1
```

```
1 x[c(0, -1)]
```

```
[1] 4 7
```

```
1 y[c(0, -1)]
```

```
[[1]]
```

```
[1] 4
```

```
[[2]]
```

```
[1] 7
```

Character subsetting

If the vector has names, selects elements whose names correspond to the values in the name vector.

```
1 x = c(a=1, b=4, c=7)
```

```
1 x["a"]
```

a

1

```
1 x[c("a", "a")]
```

a a

1 1

```
1 x[c("b", "c")]
```

b c

4 7

```
1 y = list(a=1, b=4, c=7)
```

```
1 str(y["a"])
```

List of 1

\$ a: num 1

```
1 str(y[c("a", "a")])
```

List of 2

\$ a: num 1

\$ a: num 1

```
1 str(y[c("b", "c")])
```

List of 2

\$ b: num 4

\$ c: num 7

Out of bounds

```
1 x = c(1,4,7)
```

```
1 x[4]
```

```
[1] NA
```

```
1 x[-4]
```

```
[1] 1 4 7
```

```
1 x["a"]
```

```
[1] NA
```

```
1 x[c(1,4)]
```

```
[1] 1 NA
```

```
1 y = list(1,4,7)
```

```
1 str(y[4])
```

List of 1

\$: NULL

```
1 str(y[-4])
```

List of 3

\$: num 1

\$: num 4

\$: num 7

```
1 str(y["a"])
```

List of 1

\$: NULL

```
1 str(y[c(1,4)])
```

List of 2

\$: num 1

\$: NULL

Missing values

```
1 x = c(1,4,7)
```

```
1 x[NA]
```

```
[1] NA NA NA
```

```
1 x[c(1,NA)]
```

```
[1] 1 NA
```

```
1 y = list(1,4,7)
```

```
1 str(y[NA])
```

List of 3

\$: NULL

\$: NULL

\$: NULL

```
1 str(y[c(1,NA)])
```

List of 2

\$: num 1

\$: NULL

NULL and empty vectors (length 0)

This final type of subsetting follows the rules for length coercion with a 0-length vector (i.e. the vector being subset gets coerced to having length 0 if the subsetting vector has length 0)

```
1 x = c(1,4,7)
```

```
1 x[NULL]
```

```
numeric(0)
```

```
1 x[integer()]
```

```
numeric(0)
```

```
1 x[character()]
```

```
numeric(0)
```

```
1 y = list(1,4,7)
```

```
1 y[NULL]
```

```
list()
```

```
1 y[integer()]
```

```
list()
```

```
1 y[character()]
```

```
list()
```

The other subset operators ([[and \$)

Atomic vectors - [vs. [[

[[subsets like [except it can only subset for a *single* value

```
1 x = c(a=1, b=4, c=7)
```

```
1 x[1]
```

a

1

```
1 x[[1]]
```

```
[1] 1
```

```
1 x[["a"]]
```

```
[1] 1
```

```
1 x[[1:2]]
```

Error in x[[1:2]]: attempt to select more than one element in vectorIndex

```
1 x[[TRUE]]
```

```
[1] 1
```

Generic Vectors (lists) - [vs. [[

Subsets a single value, but returns the value - not a list containing that value.
Vectors are interpreted as nested subsetting.

```
1 y = list(a=1, b=4, c=7:9)
```

```
1 y[2]
```

\$b

```
[1] 4
```

```
1 str( y[2] )
```

List of 1

```
$ b: num 4
```

```
1 y[[2]]
```

```
[1] 4
```

```
1 y[["b"]]
```

```
[1] 4
```

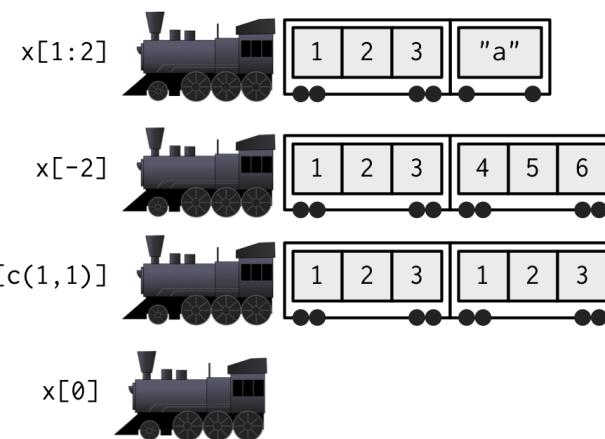
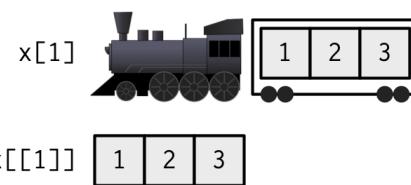
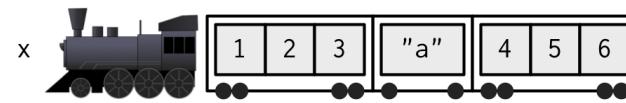
```
1 y[[1:2]]
```

Error in y[[1:2]]: subscript out of bounds

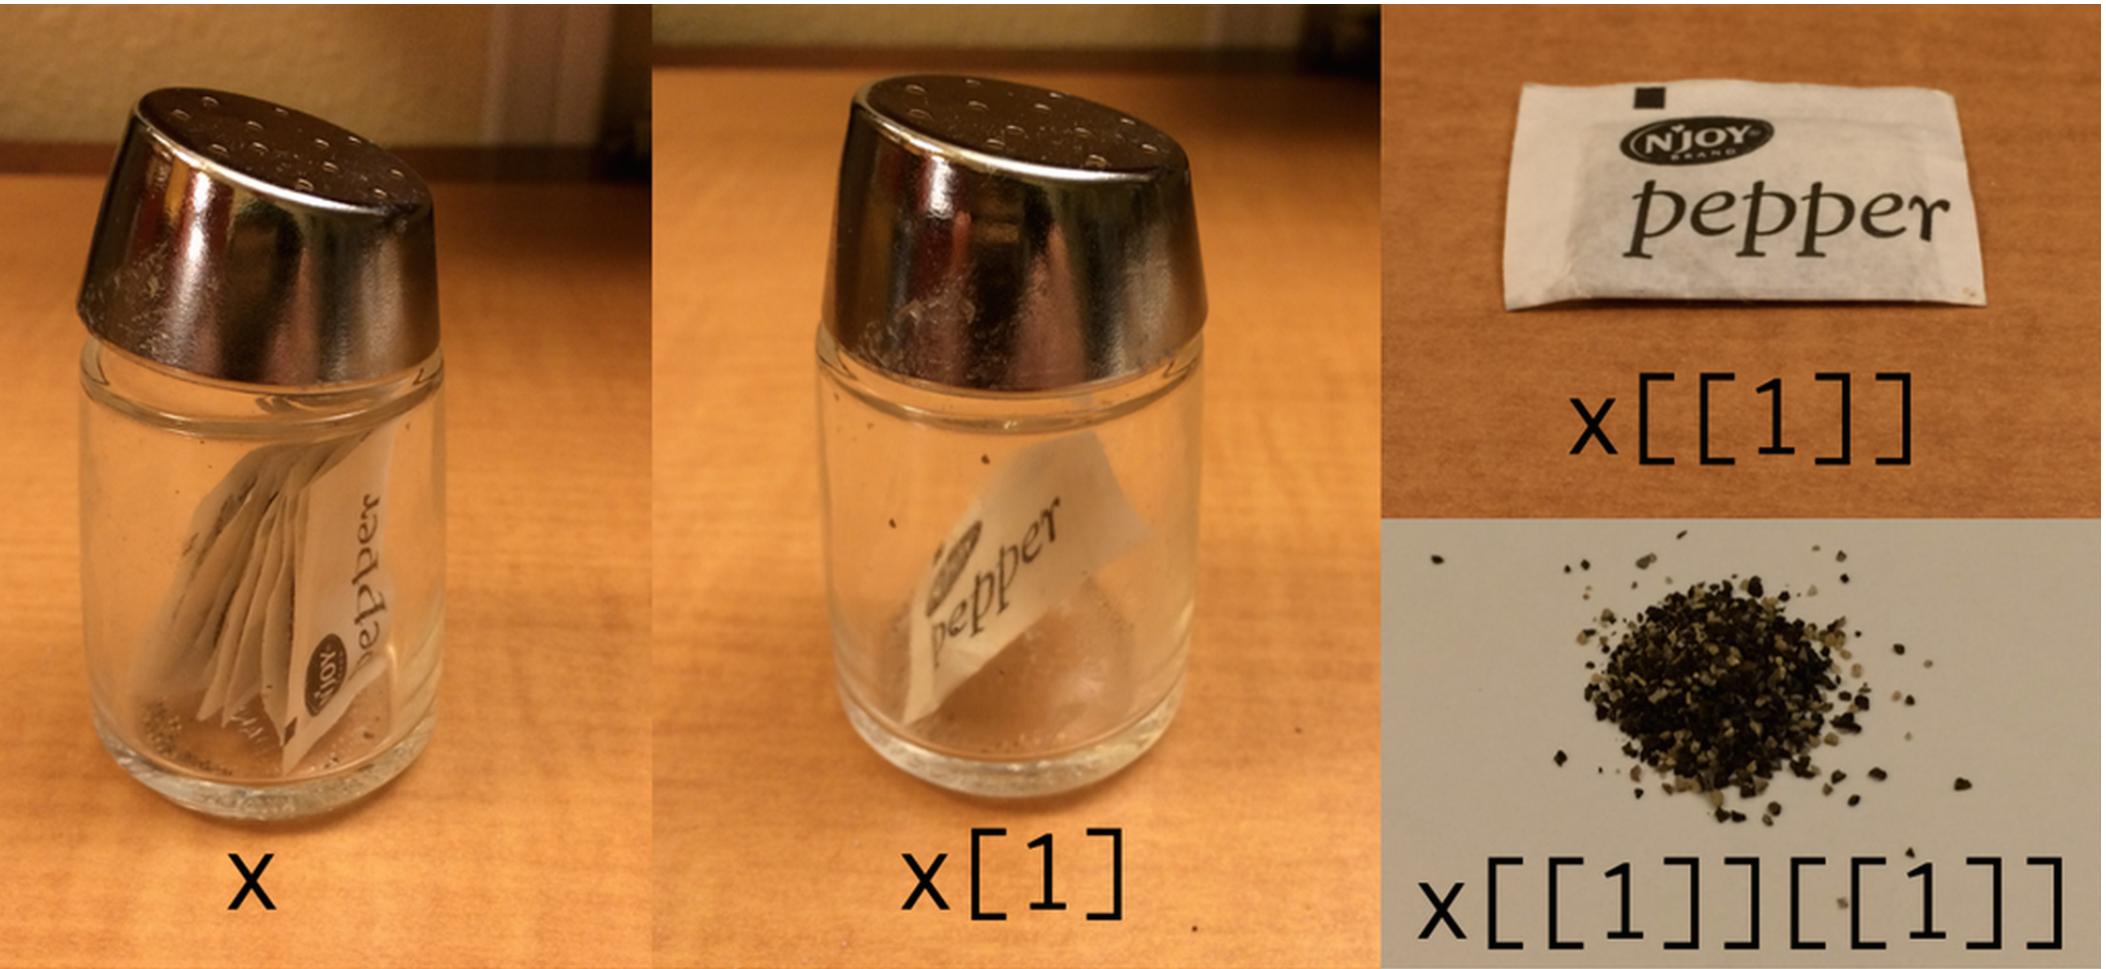
```
1 y[[2:1]]
```

```
[1] 4
```

Hadley's Analogy (1)



Hadley's Analogy (2)



Hadley Wickham @hadleywickham · 6h

Indexing lists in #rstats. Inspired by the Residence Inn



273

370

...

[[vs. \$

\$ is equivalent to [[but it only works for *named lists* and it uses partial matching for names

```
1 x = c("abc"=1, "def"=5)
```

```
1 x$abc
```

Error in x\$abc: \$ operator is invalid for atomic vectors

```
1 y = list("abc"=1, "def"=5)
```

```
1 y[["abc"]]
```

```
[1] 1
```

```
1 y$abc
```

```
[1] 1
```

```
1 y$d
```

```
[1] 5
```

A common error

Why does the following code not work?

```
1 x = list(abc = 1:10, def = 10:1)
2 y = "abc"
```

```
1 x[[y]]
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

```
1 x$y
```

```
NULL
```

The expression `x$y` gets interpreted as `x[["y"]]` by R, note the inclusion of the "s, this is not the same as the expression `x[[y]]`.

Exercise 1

Below are 100 values,

```
1 x = c(56, 3, 17, 2, 4, 9, 6, 5, 19, 5, 2, 3, 5, 0, 13, 12, 6, 31, 10, 21, 8,  
2     3, 4, 8, 5, 2, 8, 6, 18, 40, 10, 20, 1, 27, 2, 11, 14, 5, 7, 0, 3, 0, 7,  
3     21, 3, 34, 55, 18, 2, 9, 29, 1, 4, 7, 14, 7, 1, 2, 7, 4, 74, 5, 0, 3, 1,  
4     5, 2, 4, 4, 14, 15, 4, 17, 1, 9)
```

write down how you would create a subset to accomplish each of the following:

- Select every third value starting at position 2 in `x`.
- Remove all values with an odd index (e.g. 1, 3, etc.)
- Remove every 4th value, but only if it is odd.

Subsetting Data Frames

Subsetting rows

As data frames have 2 dimensions, we can subset on either the rows or the columns - the subsetting values are separated by a comma.

```
1 df = data.frame(x = 1:3, y = c("A", "B", "C"), z = TRUE))
```

```
x y z
1 1 A TRUE
2 2 B TRUE
3 3 C TRUE
```

```
1 df[1, ]
```

```
x y z
1 1 A TRUE
```

```
1 str(df[1, ])
```

```
'data.frame': 1 obs. of 3 variables:
 $ x: int 1
 $ y: chr "A"
 $ z: logi TRUE
```

```
1 df[c(1,3), ]
```

```
x y z
1 1 A TRUE
3 3 C TRUE
```

```
1 str(df[c(1,3), ])
```

```
'data.frame': 2 obs. of 3 variables:
 $ x: int 1 3
 $ y: chr "A" "C"
 $ z: logi TRUE TRUE
```

Subsetting Columns

```
1 df
```

```
x y z
1 1 A TRUE
2 2 B TRUE
3 3 C TRUE
```

```
1 df[, 1]
```

```
[1] 1 2 3
```

```
1 df[, 1:2]
```

```
x y
1 1 A
2 2 B
3 3 C
```

```
1 df[, -3]
```

```
x y
1 1 A
2 2 B
3 3 C
```

```
1 str( df[, 1] )
```

```
int [1:3] 1 2 3
```

```
1 str( df[, 1:2] )
```

```
'data.frame': 3 obs. of 2 variables:
$ x: int 1 2 3
$ y: chr "A" "B" "C"
```

```
1 str( df[, -3] )
```

```
'data.frame': 3 obs. of 2 variables:
$ x: int 1 2 3
$ y: chr "A" "B" "C"
```

Subsetting both

```
1 df
```

```
x y z
1 1 A TRUE
2 2 B TRUE
3 3 C TRUE
```

```
1 df[1, 1]
```

```
[1] 1
```

```
1 df[1:2, 1:2]
```

```
x y
1 1 A
2 2 B
```

```
1 df[-1, 2:3]
```

```
y z
2 B TRUE
3 C TRUE
```

```
1 str( df[1, 1] )
```

```
int 1
```

```
1 str( df[1:2, 1:2] )
```

```
'data.frame': 2 obs. of 2 variables:
$ x: int 1 2
$ y: chr "A" "B"
```

```
1 str( df[-1, 2:3] )
```

```
'data.frame': 2 obs. of 2
variables:
$ y: chr "B" "C"
$ z: logi TRUE TRUE
```

Preserving vs Simplifying

Most of the time, R's `[` subset operator is a *preserving* operator, in that the returned object will always have the same type/class as the object being subset.

Confusingly, when used with some classes (e.g. data frame, matrix or array) `[` becomes a *simplifying* operator (does not preserve type) - this behavior is instead controlled by the `drop` argument.

Drop

```
1 df[1, ]
```

```
x y z  
1 1 A TRUE
```

```
1 df[1, , drop=TRUE]
```

```
$x  
[1] 1  
  
$y  
[1] "A"  
  
$z  
[1] TRUE
```

```
1 df[, 1]
```

```
[1] 1 2 3
```

```
1 str(df[1, ])
```

```
'data.frame': 1 obs. of 3 variables:  
 $ x: int 1  
 $ y: chr "A"  
 $ z: logi TRUE
```

```
1 str(df[1, , drop=TRUE])
```

```
List of 3  
 $ x: int 1  
 $ y: chr "A"  
 $ z: logi TRUE
```

```
1 str(df[, 1])
```

```
int [1:3] 1 2 3
```

Exceptions

`drop` only works when the resulting value can be represented as a 1d vector (either a list or atomic).

```
1 df[1:2, 1:2]
```

```
x y  
1 1 A  
2 2 B
```

```
1 df[1:2, 1:2, drop=TRUE]
```

```
x y  
1 1 A  
2 2 B
```

```
1 str(df[1:2, 1:2])
```

```
'data.frame': 2 obs. of 2  
variables:  
$ x: int 1 2  
$ y: chr "A" "B"
```

```
1 str(df[1:2, 1:2, drop=TRUE])
```

```
'data.frame': 2 obs. of 2  
variables:  
$ x: int 1 2  
$ y: chr "A" "B"
```

Preserving vs Simplifying Subsets

Type	Simplifying	Preserving
Atomic	<code>x[[1]]</code>	<code>x[1]</code>
Vector		
List	<code>x[[1]]</code>	<code>x[1]</code>
Matrix /	<code>x[[1]]</code>	<code>x[1, , drop=FALSE]</code>
Array	<code>x[1,]</code> <code>x[, 1]</code>	<code>x[, 1, drop=FALSE]</code>
Factor	<code>x[1:4, drop=TRUE]</code>	<code>x[1:4]</code> <code>x[[1]]</code>
Data frame	<code>x[, 1]</code> <code>x[[1]]</code>	<code>x[, 1, drop=FALSE]</code> <code>x[1]</code>